



25 Recipes for

Getting Started with R

O'REILLY®

Paul Teetor

25 Recipes for Getting Started with R

25 Recipes for Getting Started with R

Paul Teetor

O'REILLY®

Beijing • Cambridge • Farnham • Köln • Sebastopol • Tokyo

25 Recipes for Getting Started with R

by Paul Teetor

Copyright © 2011 Paul Teetor. All rights reserved.

Printed in the United States of America.

Published by O'Reilly Media, Inc., 1005 Gravenstein Highway North, Sebastopol, CA 95472.

O'Reilly books may be purchased for educational, business, or sales promotional use. Online editions are also available for most titles (<http://my.safaribooksonline.com>). For more information, contact our corporate/institutional sales department: (800) 998-9938 or corporate@oreilly.com.

Editor: Mike Loukides

Production Editor: Adam Zarella

Proofreader: Adam Zarella

Cover Designer: Karen Montgomery

Interior Designer: David Futato

Illustrator: Robert Romano

Printing History:

February 2011: First Edition.

Nutshell Handbook, the Nutshell Handbook logo, and the O'Reilly logo are registered trademarks of O'Reilly Media, Inc. *25 Recipes for Getting Started with R*, the image of a harpy eagle, and related trade dress are trademarks of O'Reilly Media, Inc.

Many of the designations used by manufacturers and sellers to distinguish their products are claimed as trademarks. Where those designations appear in this book, and O'Reilly Media, Inc. was aware of a trademark claim, the designations have been printed in caps or initial caps.

While every precaution has been taken in the preparation of this book, the publisher and authors assume no responsibility for errors or omissions, or for damages resulting from the use of the information contained herein.

ISBN: 978-1-449-30323-5

[LSI]

1296229980

Table of Contents

Preface	vii
The Recipes	1
1.1 Downloading and Installing R	1
1.2 Getting Help on a Function	3
1.3 Viewing the Supplied Documentation	4
1.4 Searching the Web for Help	6
1.5 Reading Tabular Datafiles	8
1.6 Reading from CSV Files	10
1.7 Creating a Vector	12
1.8 Computing Basic Statistics	13
1.9 Initializing a Data Frame from Column Data	16
1.10 Selecting Data Frame Columns by Position	17
1.11 Selecting Data Frame Columns by Name	21
1.12 Forming a Confidence Interval for a Mean	22
1.13 Forming a Confidence Interval for a Proportion	23
1.14 Comparing the Means of Two Samples	24
1.15 Testing a Correlation for Significance	26
1.16 Creating a Scatter Plot	28
1.17 Creating a Bar Chart	29
1.18 Creating a Box Plot	30
1.19 Creating a Histogram	32
1.20 Performing Simple Linear Regression	33
1.21 Performing Multiple Linear Regression	34
1.22 Getting Regression Statistics	36
1.23 Diagnosing a Linear Regression	39
1.24 Predicting New Values	42
1.25 Accessing the Functions in a Package	43

Preface

R is a powerful tool for statistics, graphics, and statistical programming. It is used by tens of thousands of people daily to perform serious statistical analyses. It is a free, open source system whose implementation is the collective accomplishment of many intelligent, hard-working people. There are more than 2,000 available add-ons, and R is a serious rival to all commercial statistical packages.

But R can be frustrating. It's not obvious how to accomplish many tasks, even simple ones. The simple tasks are easy once you know how, yet figuring out the "how" can be maddening.

This is a book of how-to recipes for beginners, each of which solves a specific problem. The recipe includes a quick introduction to the solution, followed by a discussion that aims to unpack the solution and give you some insight into how it works. I know these recipes are useful and I know they work because I use them myself.

Most recipes use one or two R functions to solve the stated problem. It's important to remember that I do not describe the functions in detail; rather, I describe just enough to get the job done. Nearly every such function has additional capabilities beyond those described here, and some of those capabilities are amazing. I strongly urge you to read the function's help page. You will likely learn something valuable.

The book is not a tutorial on R, although you will learn something by studying the recipes. The book is not an introduction to statistics, either. The recipes assume that you are familiar with the underlying statistical procedure, if any, and just want to know how it's done in R.

These recipes were taken from my [R Cookbook](#) (O'Reilly). The *Cookbook* contains over 200 recipes that you will find useful when you move beyond the basics of R.

Other Resources

I can recommend several other resources for R beginners:

An Introduction to R (Network Theory Limited)

This book by William N. Venables, et al., covers many general topics, including statistics, graphics, and programming. You can download the [free PDF book](#); or, better yet, buy the printed copy because the profits are donated to the R project.

R in a Nutshell (O'Reilly)

Joseph Adler's book is the tutorial and reference you'll keep by your side. It covers many topics, from introductory material to advanced techniques.

Using R for Introductory Statistics (Chapman & Hall/CRC)

A good choice for learning R and statistics together by John Verzani. The book teaches statistical concepts together with the skills needed to apply them using R.

The R community has also produced many tutorials and introductions, especially in specialized topics. Most of this material is available on the Web, so I suggest searching there when you have a specific need (as in [Recipe 1.4](#)).

The R project website keeps an extensive bibliography of [books related to R](#), both for beginning and advanced users.

Downloading Additional Packages

The R project has over 2,000 packages that you can download to augment the standard distribution with additional capabilities. You might see such packages mentioned in the See Also section of a recipe, or you might discover one while searching the Web.

Most packages are available through the Comprehensive R Archive Network (CRAN) at <http://cran.r-project.org>. From the CRAN home page, click on Packages to see the name and a brief description of every available package. Click on a package name to see more information, including the package documentation.

Downloading and installing a package is simple via the `install.packages` function. You would install the `zoo` package this way, for example:

```
> install.packages("zoo")
```

When R prompts you for a mirror site, select one near you. R will download both the package and any packages on which it depends, then install them onto your machine.

On Linux or Unix, I suggest having the systems administrator install packages into the system-wide directories, making them available to all users. If that is not possible, install the packages into your private directories.

Software and Platform Notes

The base distribution of R has frequent, planned releases, but the language definition and core implementation are stable. The recipes in this book should work with any recent release of the base distribution.

One recipe has platform-specific considerations ([Recipe 1.1](#)). As far as I know, all other recipes will work on all three major platforms for R: Windows, OS X, and Linux/Unix.

Conventions Used in This Book

The following typographical conventions are used in this book:

Italic

Indicates new terms, URLs, email addresses, filenames, and file extensions.

Constant width

Used for program listings, as well as within paragraphs to refer to program elements such as variable or function names, databases, data types, environment variables, statements, and keywords.

Constant width bold

Shows commands or other text that should be typed literally by the user.

Constant width italic

Shows text that should be replaced with user-supplied values or by values determined by context.



This icon signifies a tip, suggestion, or general note.



This icon indicates a warning or caution.

Using Code Examples

This book is here to help you get your job done. In general, you may use the code in this book in your programs and documentation. You do not need to contact us for permission unless you're reproducing a significant portion of the code. For example, writing a program that uses several chunks of code from this book does not require permission. Selling or distributing a CD-ROM of examples from O'Reilly books does require permission. Answering a question by citing this book and quoting example

code does not require permission. Incorporating a significant amount of example code from this book into your product's documentation does require permission.

We appreciate, but do not require, attribution. An attribution usually includes the title, author, publisher, and ISBN. For example: “*25 Recipes for Getting Started with R* by Paul Teetor (O’Reilly). Copyright 2011 Paul Teetor, 978-1-449-30323-5.”

If you feel your use of code examples falls outside fair use or the permission given above, feel free to contact us at permissions@oreilly.com.

Safari® Books Online



Safari Books Online is an on-demand digital library that lets you easily search over 7,500 technology and creative reference books and videos to find the answers you need quickly.

With a subscription, you can read any page and watch any video from our library online. Read books on your cell phone and mobile devices. Access new titles before they are available for print, and get exclusive access to manuscripts in development and post feedback for the authors. Copy and paste code samples, organize your favorites, download chapters, bookmark key sections, create notes, print out pages, and benefit from tons of other time-saving features.

O’Reilly Media has uploaded this book to the Safari Books Online service. To have full digital access to this book and others on similar topics from O’Reilly and other publishers, sign up for free at <http://my.safaribooksonline.com>.

How to Contact Us

Please address comments and questions concerning this book to the publisher:

O’Reilly Media, Inc.
1005 Gravenstein Highway North
Sebastopol, CA 95472
800-998-9938 (in the United States or Canada)
707-829-0515 (international or local)
707-829-0104 (fax)

We have a web page for this book, where we list errata, examples, and any additional information. You can access this page at:

<http://www.oreilly.com/catalog/9781449303235>

To comment or ask technical questions about this book, send email to:

bookquestions@oreilly.com

For more information about our books, courses, conferences, and news, see our website at <http://oreilly.com>.

Find us on Facebook: <http://facebook.com/oreilly>

Follow us on Twitter: <http://twitter.com/oreillymedia>

Watch us on YouTube: <http://www.youtube.com/oreillymedia>

The Recipes

1.1 Downloading and Installing R

Problem

You want to install R on your computer.

Solution

Windows and OS X users can download R from CRAN, the Comprehensive R Archive Network. Linux and Unix users can install R packages using their package management tool.

Windows

- Open <http://www.r-project.org/> in your browser.
- Click on “CRAN”. You’ll see a list of mirror sites, organized by country.
- Select a site near you.
- Click on “Windows” under “Download and Install R”.
- Click on “base”.
- Click on the link to download the latest version of R (an *.exe* file).
- When the download completes, double-click the *.exe* file and answer the usual questions.

OS X

- Open <http://www.r-project.org/> in your browser.
- Click on “CRAN”. You’ll see a list of mirror sites, organized by country.
- Select a site near you.
- Click on “MacOS X”.
- Click on the *.pkg* file for the latest version of R, under “Files:”, to download it.
- When the download completes, double-click the *.pkg* file and answer the usual questions.

Linux or Unix

The major Linux distributions have packages for installing R. Here are some examples:

Distribution	Package name
Ubuntu or Debian	r-base
Red Hat or Fedora	R.i386
Suse	R-base

Use the system's package manager to download and install the package. Normally, you will need the root password or `sudo` privileges; otherwise, ask a system administrator to perform the installation.

Discussion

Installing R on Windows or OS X is straightforward because there are prebuilt binaries for those platforms. You need only follow the preceding instructions. The CRAN Web pages also contain links to installation-related resources, such as frequently asked questions (FAQs) and tips for special situations (“How do I install R when using Windows Vista?”) that you may find useful.

Theoretically, you can install R on Linux or Unix in one of two ways: by installing a distribution package or by building it from scratch. In practice, installing a package is the preferred route. The distribution packages greatly streamline both the initial installation and subsequent updates.

On Ubuntu or Debian, use `apt-get` to download and install R. Run under `sudo` to have the necessary privileges:

```
$ sudo apt-get install r-base
```

On Red Hat or Fedora, use `yum`:

```
$ sudo yum install R.i386
```

Most platforms also have graphical package managers, which you might find more convenient.

Beyond the base packages, I recommend installing the documentation packages, too. On my Ubuntu machine, for example, I installed `r-base-html` (because I like browsing the hyperlinked documentation) as well as `r-doc-html`, which installs the important R manuals locally:

```
$ sudo apt-get install r-base-html r-doc-html
```

Some Linux repositories also include prebuilt copies of R packages available on CRAN. I don't use them because I'd rather get my software directly from CRAN itself, which usually has the freshest versions.

In rare cases, you may need to build R from scratch. You might have an obscure, unsupported version of Unix; or you might have special considerations regarding performance or configuration. The build procedure on Linux or Unix is quite standard. Download the tarball from the home page of your CRAN mirror; it's called something like *R-2.12.1.tar.gz*, except the "2.12.1" will be replaced by the latest version. Unpack the tarball, look for a file called *INSTALL*, and follow the directions.

See Also

R in a Nutshell contains more details for downloading and installing R, including instructions for building the Windows and OS X versions. Perhaps the ultimate guide, though, is *R Installation and Administration* (<http://cran.r-project.org/doc/manuals/R-admin.html>), available on CRAN, which describes how to build and install R on a variety of platforms.

This recipe is about installing the base package. Use the `install.packages` function to install add-on packages from CRAN.

1.2 Getting Help on a Function

Problem

You want to know more about a function that is installed on your machine.

Solution

Use `help` to display the documentation for the function:

```
> help(functionname)
```

Use `args` for a quick reminder of the function arguments:

```
> args(functionname)
```

Use `example` to see examples of using the function:

```
> example(functionname)
```

Discussion

I present many R functions in this book. Every R function has more bells and whistles than I can possibly describe. If a function catches your interest, I strongly suggest reading the help page for that function. One of its bells or whistles might be very useful to you.

Suppose you want to know more about the `mean` function. Use the `help` function like this:

```
> help(mean)
```


This will either open a window with function documentation or display the documentation on your console, depending on your platform. A shortcut for the `help` command is to simply type `?`, followed by the function name:

```
> ?mean
```

Sometimes you just want a quick reminder of the arguments to a function; what are they, and in what order do they occur? Use the `args` function:

```
> args(mean)
function (x, ...)
NULL
> args(sd)
function (x, na.rm = FALSE)
NULL
```

The first line of output from `args` is a synopsis of the function call. For `mean`, the synopsis shows one argument, `x`, which is a vector of numbers. For `sd`, the synopsis shows the same vector, `x`, and an optional argument called `na.rm`. (You can ignore the second line of output, which is often just `NULL`.)

Most documentation for functions includes examples near the end. A cool feature of R is that you can request that it execute the examples, giving you a little demonstration of the function's capabilities. The documentation for the `mean` function, for instance, contains examples, but you don't need to type them yourself. Just use the `example` function to watch them run:

```
> example(mean)

mean> x <- c(0:10, 50)

mean> xm <- mean(x)

mean> c(xm, mean(x, trim = 0.1))
[1] 8.75 5.50

mean> mean(USArrests, trim = 0.2)
      Murder  Assault UrbanPop      Rape
       7.42   167.60    66.20    20.16
```

The user typed `example(mean)`. Everything else was produced by R, which executed the examples from the help page and displayed the results.

1.3 Viewing the Supplied Documentation

Problem

You want to read the documentation supplied with R.

Solution

Use the `help.start` function to see the documentation's table of contents:

```
> help.start()
```

From there, links are available to all the installed documentation.

Discussion

The base distribution of R includes a wealth of documentation—literally thousands of pages. When you install additional packages, those packages contain documentation that is also installed on your machine.

It is easy to browse this documentation via the `help.start` function, which opens a window on the top-level table of contents; see [Figure 1-1](#).




Figure 1-1. Documentation table of contents

The two links in the Reference section are especially useful:

Packages

Click here to see a list of all the installed packages, both in the base packages and the additional installed packages. Click on a package name to see a list of its functions and datasets.

Search Engine & Keywords

Click here to access a simple search engine, which allows you to search the documentation by keyword or phrase. There is also a list of common keywords, organized by topic; click one to see the associated pages.

See Also

The local documentation is copied from the [R Project website](#), which may have updated documents.

1.4 Searching the Web for Help

Problem

You want to search the Web for information and answers regarding R.

Solution

Inside R, use the `RSiteSearch` function to search by keyword or phrase:

```
> RSiteSearch("key phrase")
```

In your browser, try using these sites for searching:

<http://rseek.org>

This is a Google custom search that is focused on R-specific websites.

<http://stackoverflow.com/>

Stack Overflow is a searchable Q&A site oriented toward programming issues such as data structures, coding, and graphics.

<http://stats.stackexchange.com/>

The Statistical Analysis area on Stack Exchange is also a searchable Q&A site, but it is oriented more toward statistics than programming.

Discussion

The `RSiteSearch` function will open a browser window and direct it to the search engine on the R Project website (<http://search.r-project.org/>). There, you will see an initial search that you can refine. For example, this call would start a search for “canonical correlation”:

```
> RSiteSearch("canonical correlation")
```

This is quite handy for doing quick Web searches without leaving R. However, the search scope is limited to R documentation and the mailing-list archives.

RSeek.org provides a wider search. Its virtue is that it harnesses the power of the Google search engine while focusing on sites relevant to R. That eliminates the extraneous results of a generic Google search. The beauty of RSeek.org is that it organizes the results in a useful way.

Figure 1-2 shows the results of visiting RSeek.org and searching for “canonical correlation”. The left side of the page shows general results for search R sites. The right side is a tabbed display that organizes the search results into several categories:

- Introductions
- Task Views
- Support Lists
- Functions
- Books
- Blogs
- Related Tools

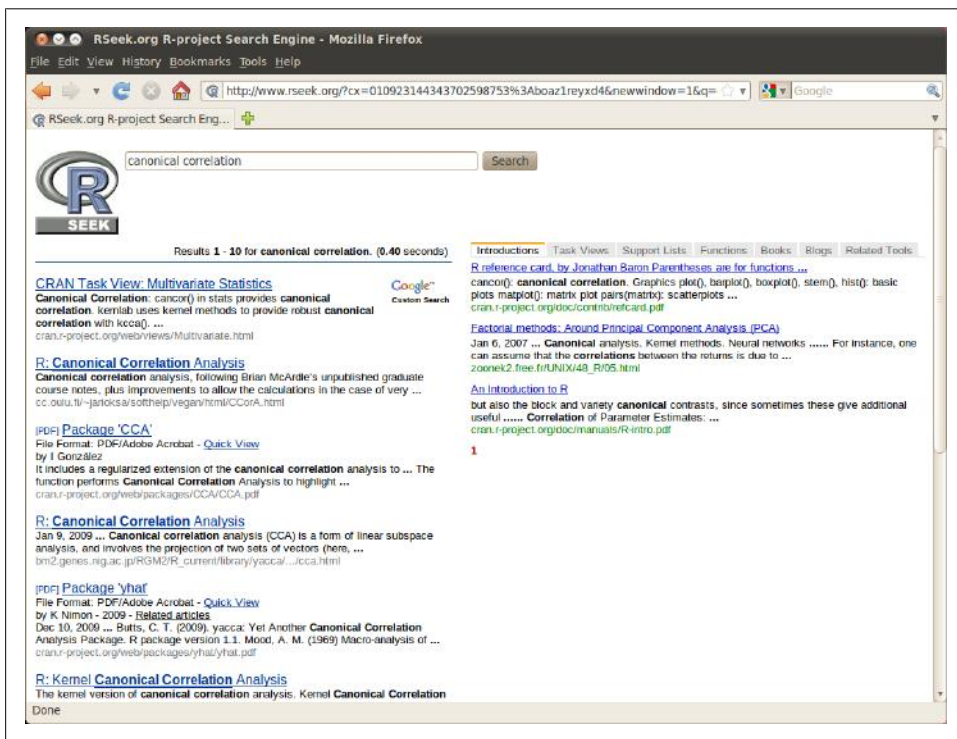


Figure 1-2. Search results from RSeek.org

If you click on the Introductions tab, for example, you'll find tutorial material. The Task Views tab will show any Task View that mentions your search term. Likewise, clicking on Functions will show links to relevant R functions. This is a good way to zero in on search results.

[Stack Overflow](#) is a so-called Q&A site, which means that anyone can submit a question and experienced users will supply answers—often there are multiple answers to each question. Readers vote on the answers, so good answers tend to rise to the top. This creates a rich database of searchable Q&A dialogs. Stack Overflow is strongly problem oriented, and the topics lean toward the programming side of R.

Stack Overflow hosts questions for many programming languages; therefore, when entering a term into their search box, prefix it with “[r]” to focus the search on questions tagged for R. For example, searching for “[r] standard error” will select only the questions tagged for R and will avoid Python and C++ questions.

Stack Exchange (not Overflow) has a Q&A area for Statistical Analysis. The area is more focused on statistics than programming, so use this site when seeking answers that are more concerned with statistics in general and less with R in particular.

See Also

If your search reveals a useful package, use the `install.packages` function to install it on your machine.

1.5 Reading Tabular Datafiles

Problem

You want to read a text file that contains a table of data.

Solution

Use the `read.table` function, which returns a data frame:

```
> dfm <- read.table("filename.txt")
```

Discussion

Tabular datafiles are quite common. They are text files with a simple format:

- Each line contains one record.
- Within each record, fields (items) are separated by a one-character delimiter, such as a space, tab, colon, or comma.
- Each record contains the same number of fields.

This format is more free-form than the fixed-width format because fields needn't be aligned by position. Here is a datafile in tabular format, called `statisticians.txt`, using a space character between fields:

```
Fisher R.A. 1890 1962
Pearson Karl 1857 1936
Cox Gertrude 1900 1978
Yates Frank 1902 1994
Smith Kirstine 1878 1939
```

The `read.table` function is built to read this file. By default, it assumes the data fields are separated by white space (blanks or tabs):

```
> dfrm <- read.table("statisticians.txt")
> print(dfrm)
      V1      V2  V3  V4
1 Fisher   R.A. 1890 1962
2 Pearson  Karl 1857 1936
3   Cox Gertrude 1900 1978
4   Yates  Frank 1902 1994
5  Smith Kirstine 1878 1939
```

If your file uses a separator other than white space, specify it using the `sep` parameter. For example, if our file used a colon (:) as the field separator, we would read it this way:

```
> dfrm <- read.table("statisticians.txt", sep=":")
```

You can't tell from the printed output, but `read.table` interpreted the first and last names as factors, not strings. We see that by checking the class of the resulting column:

```
> class(dfrm$V1)
[1] "factor"
```

To prevent `read.table` from interpreting character strings as factors, set the `stringsAsFactors` parameter to `FALSE`:

```
> dfrm <- read.table("statisticians.txt", stringsAsFactor=FALSE)
> class(dfrm$V1)
[1] "character"
```

Now the class of the first column is character, not factor.

If any field contains the string "NA", then `read.table` assumes that the value is missing and converts it to NA. Your datafile might employ a different string to signal missing values; if it does, use the `na.strings` parameter. The SAS convention, for example, is that missing values are signaled by a single period (.). We can read such datafiles like this:

```
> dfrm <- read.table("filename.txt", na.strings=".")
```

I am a huge fan of *self-describing data*: datafiles which describe their own contents. (A computer scientist would say the file contains its own *metadata*.) The `read.table` function has two features that support this characteristic. First, you can include a *header line* at the top of your file that gives names to the columns. The line contains one name

for every column, and it uses the same field separator as the data. Here is our datafile with a header line that names the columns:

```
lastname firstname born died
Fisher R.A. 1890 1962
Pearson Karl 1857 1936
Cox Gertrude 1900 1978
Yates Frank 1902 1994
Smith Kirstine 1878 1939
```

Now we can tell `read.table` that our file contains a header line, and it will use the column names when it builds the data frame:

```
> dfrm <- read.table("statisticians.txt", header=TRUE, stringsAsFactor=FALSE)
> print(dfrm)
  lastname firstname born died
1  Fisher      R.A. 1890 1962
2  Pearson     Karl 1857 1936
3    Cox  Gertrude 1900 1978
4   Yates     Frank 1902 1994
5   Smith  Kirstine 1878 1939
```

The second feature of `read.table` is *comment lines*. Any line that begins with a pound sign (#) is ignored, so you can put comments on those lines:

```
# This is a datafile of famous statisticians.
# Last edited on 1994-06-18
lastname firstname born died
Fisher R.A. 1890 1962
Pearson Karl 1857 1936
Cox Gertrude 1900 1978
Yates Frank 1902 1994
Smith Kirstine 1878 1939
```

`read.table` has many parameters for controlling how it reads and interprets the input file. See the help page for details.

See Also

If your data items are separated by commas, see [Recipe 1.6](#) for reading a CSV file.

1.6 Reading from CSV Files

Problem

You want to read data from a comma-separated values (CSV) file.

Solution

The `read.csv` function can read CSV files. If your CSV file has a header line, use this:

```
> tbl <- read.csv("filename")
```

If your CSV file does not contain a header line, set the `header` option to `FALSE`:

```
> tbl <- read.csv("filename", header=FALSE)
```

Discussion

The CSV file format is popular because many programs can import and export data in that format. Such programs include R, Excel, other spreadsheet programs, many database managers, and most statistical packages. CSV is a flat file of tabular data, in which each line in the file is a row of data, and each row contains data items separated by commas. Here is a very simple CSV file with three rows and three columns (the first line is a *header line* that contains the column names, also separated by commas):

```
label,lbound,ubound
low,0,0.674
mid,0.674,1.64
high,1.64,2.33
```

The `read.csv` function reads the data and creates a data frame, which is the usual R representation for tabular data. The function assumes that your file has a header line unless told otherwise:

```
> tbl <- read.csv("table-data.csv")
> tbl
  label lbound ubound
1 low  0.000  0.674
2 mid  0.674  1.640
3 high 1.640  2.330
```

Observe that `read.csv` took the column names from the header line for the data frame. If the file did not contain a header, we would specify `header=FALSE`, and R would synthesize column names for us (`V1`, `V2`, and `V3` in this case):

```
> tbl <- read.csv("table-data-with-no-header.csv", header=FALSE)
> tbl
  V1 V2 V3
1 low 0.000 0.674
2 mid 0.674 1.640
3 high 1.640 2.330
```

A good feature of `read.csv` is that it automatically interprets nonnumeric data as a factor (categorical variable), which is often what you want since this is, after all, a statistical package, not Perl. The `label` variable in the `tbl` data frame just shown is actually a factor, not a character variable. You see that by inspecting the structure of `tbl`:

```
> str(tbl)
'data.frame': 3 obs. of 3 variables:
 $ label : Factor w/ 3 levels "high","low","mid": 2 3 1
 $ lbound: num 0 0.674 1.64
 $ ubound: num 0.674 1.64 2.33
```

Sometimes, you really want your data interpreted as strings, not as factors. In that case, set the `as.is` parameter to `TRUE`; this indicates that R should not interpret nonnumeric data as a factor:


```
> tbl <- read.csv("table-data.csv", as.is=TRUE)
> str(tbl)
'data.frame':  3 obs. of  3 variables:
 $ label : chr  "low" "mid" "high"
 $ lbound: num  0 0.674 1.64
 $ ubound: num  0.674 1.64 2.33
```

Notice that the `label` variable now has character-string values and is no longer a factor.

Another useful feature is that input lines starting with a pound sign (`#`) are ignored, which lets you embed comments in your datafile. Disable this feature by specifying `comment.char=""`.

The `read.csv` function has many useful bells and whistles. These include the ability to skip leading lines in the input file, control the conversion of individual columns, fill out short rows, limit the number of lines, and control the quoting of strings. See the R help page for details.

See Also

See the R help page for `read.table`, which is the basis for `read.csv`. See the `write.csv` function for writing CSV files.

1.7 Creating a Vector

Problem

You want to create a vector.

Solution

Use the `c(...)` operator to construct a vector from given values.

Discussion

Vectors are a central component of R, not just another data structure. A vector can contain numbers, strings, or logical values, but not a mixture.

The `c(...)` operator can construct a vector from simple elements:

```
> c(1,1,2,3,5,8,13,21)
[1]  1  1  2  3  5  8 13 21
> c(1*pi, 2*pi, 3*pi, 4*pi)
[1] 3.141593 6.283185 9.424778 12.566371
> c("Everyone", "loves", "stats.")
[1] "Everyone" "loves"    "stats."
> c(TRUE, TRUE, FALSE, TRUE)
[1] TRUE TRUE FALSE TRUE
```

If the arguments to `c(...)` are themselves vectors, it flattens them and combines them into a single vector:

```
> v1 <- c(1,2,3)
> v2 <- c(4,5,6)
> c(v1,v2)
[1] 1 2 3 4 5 6
```

Vectors cannot contain a mix of data types, such as numbers and strings. If you create a vector from mixed elements, R will try to accommodate you by converting one of them:

```
> v1 <- c(1,2,3)
> v3 <- c("A","B","C")
> c(v1,v3)
[1] "1" "2" "3" "A" "B" "C"
```


Here, the user tried to create a vector from both numbers and strings. R converted all the numbers to strings before creating the vector, thereby making the data elements compatible.

Technically speaking, two data elements can coexist in a vector only if they have the same *mode*. The modes of `3.1415` and `"foo"` are numeric and character, respectively:

```
> mode(3.1415)
[1] "numeric"
> mode("foo")
[1] "character"
```

Those modes are incompatible. To make a vector from them, R converts `3.1415` to character mode so it will be compatible with `"foo"`:

```
> c(3.1415, "foo")
[1] "3.1415" "foo"
> mode(c(3.1415, "foo"))
[1] "character"
```



`c` is a generic operator, which means that it works with many datatypes and not just vectors. However, it might not do exactly what you expect, so check its behavior before applying it to other datatypes and objects.

1.8 Computing Basic Statistics

Problem

You want to calculate basic statistics: mean, median, standard deviation, variance, correlation, or covariance.

Solution

Use one of these functions as appropriate, assuming that `x` and `y` are vectors:

- `mean(x)`
- `median(x)`
- `sd(x)`
- `var(x)`
- `cor(x, y)`
- `cov(x, y)`

Discussion

When I first opened the documentation for R, I began searching for material called something like “Procedures for Calculating Standard Deviation.” I figured that such an important topic would likely require a whole chapter.

It’s not that complicated.

Standard deviation and other basic statistics are calculated by simple functions. Ordinarily, the function argument is a vector of numbers, and the function returns the calculated statistic:

```
> x <- c(0,1,1,2,3,5,8,13,21,34)
> mean(x)
[1] 8.8
> median(x)
[1] 4
> sd(x)
[1] 11.03328
> var(x)
[1] 121.7333
```

The `sd` function calculates the sample standard deviation, and `var` calculates the sample variance.

The `cor` and `cov` functions can calculate the correlation and covariance, respectively, between two vectors:

```
> x <- c(0,1,1,2,3,5,8,13,21,34)
> y <- log(x+1)
> cor(x,y)
[1] 0.9068053
> cov(x,y)
[1] 11.49988
```

All these functions are picky about values that are not available (NA). Even one NA value in the vector argument causes any of these functions to return NA, or even halt altogether with a cryptic error:

```
> x <- c(0,1,1,2,3,NA)
> mean(x)
[1] NA
> sd(x)
[1] NA
```

It's annoying when R is that cautious, but it is the right thing to do. You must think carefully about your situation. Does an NA in your data invalidate the statistic? If yes, then R is doing the right thing. If not, you can override this behavior by setting `na.rm=TRUE`, which tells R to ignore the NA values:

```
> x <- c(0,1,1,2,3,NA)
> mean(x, na.rm=TRUE)
[1] 1.4
> sd(x, na.rm=TRUE)
[1] 1.140175
```

A beautiful aspect of `mean` and `sd` is that they are smart about data frames. They understand that each column of the data frame is a different variable, so they calculate their statistic for each column individually. This example calculates those basic statistics for a data frame with three columns:

```
> print(dframe)
      small  medium    big
1  0.6739635 10.526448 99.83624
2  1.5524619  9.205156 100.70852
3  0.3250562 11.427756 99.73202
4  1.2143595  8.533180 98.53608
5  1.3107692  9.763317 100.74444
6  2.1739663  9.806662 98.58961
7  1.6187899  9.150245 100.46707
8  0.8872657 10.058465 99.88068
9  1.9170283  9.182330 100.46724
10 0.7767406  7.949692 100.49814
> mean(dframe)
      small  medium    big
1.245040  9.560325 99.946003
> sd(dframe)
      small  medium    big
0.5844025  0.9920281 0.8135498
```

Notice that `mean` and `sd` both return three values, one for each column defined by the data frame. (Technically, they return a three-element vector in which the `names` attribute is taken from the columns of the data frame.)

The `var` function understands data frames, too, but it behaves quite differently from `mean` and `sd`. It calculates the covariance between the columns of the data frame and returns the covariance matrix:

```
> var(dframe)
      small  medium    big
small  0.34152627 -0.21516416 -0.04005275
medium -0.21516416  0.98411974 -0.09253855
big    -0.04005275 -0.09253855  0.66186326
```

Likewise, if `x` is either a data frame or a matrix, `cor(x)` returns the correlation matrix and `cov(x)` returns the covariance matrix:

```
> cor(dframe)
      small      medium      big
small 1.00000000 -0.3711367 -0.08424345
medium -0.37113670 1.00000000 -0.11466070
big    -0.08424345 -0.1146607 1.00000000
> cov(dframe)
      small      medium      big
small 0.34152627 -0.21516416 -0.04005275
medium -0.21516416 0.98411974 -0.09253855
big    -0.04005275 -0.09253855 0.66186326
```

Alas, the `median` function does not understand data frames. To calculate the medians of data frame columns, use the `lapply` function to apply the `median` function to each column separately.

See Also

See [Recipe 1.12](#) for calculating the confidence interval of the mean. See [Recipe 1.15](#) for testing the significance of a correlation.

1.9 Initializing a Data Frame from Column Data

Problem

Your data is organized by columns, and you want to assemble it into a data frame.

Solution

If your data is captured in several vectors and/or factors, use the `data.frame` function to assemble them into a data frame:

```
> dfrm <- data.frame(v1, v2, v3, f1, f2)
```

Use `as.data.frame` instead if your data is captured in a *list* that contains vectors and/or factors:

```
> dfrm <- as.data.frame(list.of.vectors)
```

Discussion

A data frame is a collection of columns, each of which corresponds to an observed variable (in the statistical sense, not the programming sense). If your data is already organized into columns, then it's easy to build a data frame.

The `data.frame` function can construct a data frame from vectors, where each vector is one observed variable. Suppose you have two numeric predictor variables, one categorical predictor variable, and one response variable. The `data.frame` function can create a data frame from your vectors:

```
> dfrm <- data.frame(pred1, pred2, pred3, resp)
> dfrm
      pred1      pred2 pred3      resp
1 -2.7528917 -1.40784130 AM 12.57715
2 -0.3626909  0.31286963 AM 21.02418
3 -1.0416039 -0.69685664 PM 18.94694
4  1.2666820 -1.27511434 PM 18.98153
5  0.7806372 -0.27292745 AM 19.59455
6 -1.0832624  0.73383339 AM 20.71605
7 -2.0883305  0.96816822 PM 22.70062
8 -0.7063653 -0.84476203 PM 18.40691
9 -0.8394022  0.31530793 PM 21.00930
10 -0.4966884 -0.08030948 AM 19.31253
```

Notice that `data.frame` takes the column names from your program variables. You can override that default by supplying explicit column names:

```
> dfrm <- data.frame(p1=pred1, p2=pred2, p3=pred3, r=resp)
> dfrm
      p1      p2 p3      r
1 -2.7528917 -1.40784130 AM 12.57715
2 -0.3626909  0.31286963 AM 21.02418
3 -1.0416039 -0.69685664 PM 18.94694
.
. (etc.)
.
```

Alternatively, your data may be organized into vectors, but those vectors are held in a list instead of individual program variables, like this:

```
> lst <- list(p1=pred1, p2=pred2, p3=pred3, r=resp)
```

No problem. Use the `as.data.frame` function to create a data frame from the list of vectors:

```
> as.data.frame(lst)
      p1      p2 p3      r
1 -2.7528917 -1.40784130 AM 12.57715
2 -0.3626909  0.31286963 AM 21.02418
3 -1.0416039 -0.69685664 PM 18.94694
.
. (etc.)
.
```

1.10 Selecting Data Frame Columns by Position

Problem

You want to select columns from a data frame according to their position.

Solution

To select a single column, use this list operator:

```
dfrm[[n]]
```

Returns *one column*—specifically, the *n*th column of *dfrm*.

To select one or more columns and package them in a data frame, use the following sublist expressions:

```
dfrm[n]
```

Returns a data frame consisting solely of the *n*th column of *dfrm*.

```
dfrm[c(n1, n2, ..., nk)]
```

Returns a data frame built from the columns in positions *n*₁, *n*₂, ..., *n*_k of *dfrm*.

You can use matrix-style subscripting to select one or more columns:

```
dfrm[, n]
```

Returns the *n*th column (assuming that *n* contains exactly one value).

```
dfrm[, c(n1, n2, ..., nk)]
```

Returns a data frame built from the columns in positions *n*₁, *n*₂, ..., *n*_k.

Note that the matrix-style subscripting can return two different data types (either column or data frame) depending on whether you select one column or multiple columns.

Discussion

There are a bewildering number of ways to select columns from a data frame. The choices can be confusing until you understand the logic behind the alternatives. As you read this explanation, notice how a slight change in syntax—a comma here, a double-bracket there — changes the meaning of the expression.

Let's play with the population data for the 16 largest cities in the Chicago metropolitan area:

```
> suburbs
```

	city	county	state	pop
1	Chicago	Cook	IL	2853114
2	Kenosha	Kenosha	WI	90352
3	Aurora	Kane	IL	171782
4	Elgin	Kane	IL	94487
5	Gary	Lake(IN)	IN	102746
6	Joliet	Kendall	IL	106221
7	Naperville	DuPage	IL	147779
8	Arlington Heights	Cook	IL	76031
9	Bolingbrook	Will	IL	70834
10	Cicero	Cook	IL	72616
11	Evanston	Cook	IL	74239
12	Hammond	Lake(IN)	IN	83048
13	Palatine	Cook	IL	67232
14	Schaumburg	Cook	IL	75386
15	Skokie	Cook	IL	63348
16	Waukegan	Lake(IL)	IL	91452

Use simple list notation to select exactly one column, such as the first column:

```
> suburbs[[1]]
[1] "Chicago"      "Kenosha"      "Aurora"      "Elgin"
[5] "Gary"         "Joliet"       "Naperville"  "Arlington Heights"
[9] "Bolingbrook"  "Cicero"       "Evanston"    "Hammond"
[13] "Palatine"     "Schaumburg"   "Skokie"     "Waukegan"
```

The first column of `suburbs` is a vector, so that's what `suburbs[[1]]` returns: a vector. If the first column were a factor, we'd get a factor.

The result differs when you use the single-bracket notation, as in `suburbs[1]` or `suburbs[c(1,3)]`. You still get the requested columns, but R wraps them in a data frame. This example returns the first column wrapped in a data frame:

```
> suburbs[1]
  city
1  Chicago
2  Kenosha
3  Aurora
4  Elgin
5  Gary
6  Joliet
7  Naperville
8  Arlington Heights
9  Bolingbrook
10 Cicero
11 Evanston
12 Hammond
13 Palatine
14 Schaumburg
15 Skokie
16 Waukegan
```

The next example returns the first and third columns wrapped in a data frame:

```
> suburbs[c(1,3)]
  city      pop
1  Chicago 2853114
2  Kenosha  90352
3  Aurora  171782
4  Elgin   94487
5  Gary   102746
6  Joliet  106221
7  Naperville 147779
8  Arlington Heights 76031
9  Bolingbrook 70834
10 Cicero    72616
11 Evanston  74239
12 Hammond  83048
13 Palatine  67232
14 Schaumburg 75386
15 Skokie    63348
16 Waukegan  91452
```


A major source of confusion is that `suburbs[[1]]` and `suburbs[1]` look similar but produce very different results:

```
suburbs[[1]]
```

This returns one column.

```
suburbs[1]
```

This returns a data frame, and the data frame contains exactly one column. This is a special case of `dfrm[c(n_1, n_2, \dots, n_k)]`. We don't need the `c(...)` construct because there is only one n .

The point here is that “one column” is different from “a data frame that contains one column.” The first expression returns a column, so it's a vector or a factor. The second expression returns a data frame, which is different.

R lets you use matrix notation to select columns, as shown in the Solution. But an odd quirk can bite you: you might get a column or you might get a data frame, depending which many subscripts you use. In the simple case of one index you get a column, like this:

```
> suburbs[,1]
[1] "Chicago"      "Kenosha"      "Aurora"       "Elgin"
[5] "Gary"         "Joliet"       "Naperville"   "Arlington Heights"
[9] "Bolingbrook"  "Cicero"       "Evanston"     "Hammond"
[13] "Palatine"     "Schaumburg"   "Skokie"      "Waukegan"
```

But using the same matrix-style syntax with multiple indexes returns a data frame:

```
> suburbs[,c(1,4)]
  city      pop
1  Chicago 2853114
2  Kenosha  90352
3  Aurora  171782
4  Elgin   94487
5  Gary   102746
6  Joliet  106221
7  Naperville 147779
8  Arlington Heights 76031
9  Bolingbrook 70834
10 Cicero      72616
11 Evanston   74239
12 Hammond   83048
13 Palatine   67232
14 Schaumburg 75386
15 Skokie     63348
16 Waukegan   91452
```

This creates a problem. Suppose you see this expression in some old R script:

```
dfrm[,vec]
```

Does that return a column or a data frame? Well, it depends. If `vec` contains one value, you get a column; otherwise, you get a data frame. You cannot tell from the syntax alone.

To avoid this problem, you can include `drop=FALSE` in the subscripts; this forces R to return a data frame:

```
dfrm[,vec,drop=FALSE]
```

Now there is no ambiguity about the returned data structure. It's a data frame.

When all is said and done, using matrix notation to select columns from data frames is not the best procedure. I recommend that you instead use the list operators described previously. They just seem clearer.

1.11 Selecting Data Frame Columns by Name

Problem

You want to select columns from a data frame according to their name.

Solution

To select a single column, use one of these list expressions:

```
dfrm[["name"]]
```

Returns *one column*, the column called *name*.

```
dfrm$name
```

Returns the same, but presented in different syntax.

To select one or more columns and package them in a data frame, use these list expressions:

```
dfrm["name"]
```

Selects one column and packages it inside a data frame object.

```
dfrm[c("name1", "name2", ..., "namek")]
```

Selects several columns and packages them in a data frame.

You can use matrix-style subscripting to select one or more columns:

```
dfrm[, "name"]
```

Returns the named column.

```
dfrm[, c("name1", "name2", ..., "namek")]
```

Selects several columns and packages them in a data frame.

Once again, the matrix-style subscripting can return two different data types (column or data frame) depending on whether you select one column or multiple columns.

Discussion

All columns in a data frame must have names. If you know the name, it's usually more convenient and readable to select by name, rather than by position.

The solutions just described are similar to those for [Recipe 1.10](#), where we selected columns by position. The only difference is that here we use column names instead of column numbers. All the observations made in [Recipe 1.10](#) apply here:

- `dfrm[["name"]]` returns one column, not a data frame.
- `dfrm[c("name1", "name2", ..., "namek")]` returns a data frame, not a column.
- `dfrm["name"]` is a special case of the previous expression and so returns a data frame, not a column.
- The matrix-style subscripting can return either a column or a data frame, so be careful how many names you supply. See [Recipe 1.10](#) for a discussion of this “gotcha” and using `drop=FALSE`.

There is one new addition:

```
dfrm$name
```

This is identical in effect to `dfrm[["name"]]`, but it’s easier to type and to read.

See Also

See [Recipe 1.10](#) to understand more about selecting columns.

1.12 Forming a Confidence Interval for a Mean

Problem

You have a sample from a population. Given that sample, you want to determine a confidence interval for the population’s mean.

Solution

Apply the `t.test` function to your sample `x`:

```
> t.test(x)
```

The output includes a confidence interval at the 95% confidence level. To see intervals at other levels, use the `conf.level` argument.

If your sample size n is small, the underlying population must be normally distributed for there to be a meaningful confidence interval. A good rule of thumb is that “small” means $n < 30$.

Discussion

Applying the `t.test` function to a vector yields a lot of output. Buried in the output is a confidence interval:

```
> x <- rnorm(50, mean=100, sd=15)
> t.test(x)
```

One Sample t-test

```
data: x
t = 59.2578, df = 49, p-value < 2.2e-16
alternative hypothesis: true mean is not equal to 0
95 percent confidence interval:
 97.16167 103.98297
sample estimates:
mean of x
100.5723
```

In this example, the confidence interval is approximately $97.16 < \mu < 103.98$, which is sometimes written simply as (97.16, 103.98).

We can raise the confidence level to 99% by setting `conf.level=0.99`:

```
> t.test(x, conf.level=0.99)
```

One Sample t-test

```
data: x
t = 59.2578, df = 49, p-value < 2.2e-16
alternative hypothesis: true mean is not equal to 0
99 percent confidence interval:
 96.0239 105.1207
sample estimates:
mean of x
100.5723
```

That change widens the confidence interval to $96.02 < \mu < 105.12$.

1.13 Forming a Confidence Interval for a Proportion

Problem

You have a sample of values from a population consisting of successes and failures. Based on the sample data, you want to form a confidence interval for the population's proportion of successes.

Solution

Use the `prop.test` function. Suppose the sample size is n and the sample contains x successes:

```
> prop.test(n, x)
```

The function output includes the confidence interval for p .

Discussion

I subscribe to a stock market newsletter that is well written for the most part, but includes a section purporting to identify stocks that are likely to rise. It does this by looking for a certain pattern in the stock price. It recently reported, for example, that a certain stock was following the pattern. It also reported that the stock rose six times after the last nine times that pattern occurred. The writers concluded that the probability of the stock rising again was therefore $6/9$ or 66.7%.

Using `prop.test`, we can obtain the confidence interval for the true proportion of times the stock rises after the pattern. Here, the number of observations is $n = 9$ and the number of successes is $x = 6$. The output shows a confidence interval of (0.309, 0.910) at the 95% confidence level:

```
> prop.test(6, 9)

1-sample proportions test with continuity correction

data: 6 out of 9, null probability 0.5
X-squared = 0.4444, df = 1, p-value = 0.505
alternative hypothesis: true p is not equal to 0.5
95 percent confidence interval:
 0.3091761 0.9095817
sample estimates:
      p 
0.6666667
```

The writers are pretty foolish to say the probability of rising is 66.7%. They could be leading their readers into a very bad bet.

By default, `prop.test` calculates a confidence interval at the 95% confidence level. Use the `conf.level` argument for other confidence levels, like this:

```
> prop.test(n, x, p, conf.level=0.99) # 99% confidence level
```

1.14 Comparing the Means of Two Samples

Problem

You have one sample each from two populations. You want to know if the two populations could have the same mean.

Solution

Perform a t test by calling the `t.test` function:

```
> t.test(x, y)
```

By default, `t.test` assumes that your data is not paired. If the observations are paired (i.e., if each x_i is paired with the corresponding y_i), then specify `paired=TRUE`:

```
> t.test(x, y, paired=TRUE)
```

In either case, `t.test` will compute a p -value. Conventionally, if $p < 0.05$, the means are likely different, whereas $p > 0.05$ provides no such evidence:

- If either sample size is small, the populations must be normally distributed. Here, “small” means fewer than 20 data points.
- If the two populations have the same variance, specify `var.equal=TRUE` to obtain a less conservative test.

Discussion

I often use the t test to get a quick sense of the difference between two population means. It requires that the samples be large enough—both samples have 20 or more observations—or that the underlying populations be normally distributed. I don’t take the “normally distributed” part too literally. Being bell-shaped should be good enough.

A key distinction here is whether your data contains paired observations, since the answer may cause your results to differ. Suppose we want to know if coffee in the morning improves scores on SAT tests. We could run the experiment two ways:

- Randomly select one group of people. Give them the SAT test twice, once with morning coffee and once without morning coffee. For each person, we will have two SAT scores. These are *paired observations*.
- Randomly select two groups of people. One group has a cup of morning coffee and takes the SAT test. The other group just takes the test. We have a score for each person, but the scores are not paired in any way.

Statistically, these experiments are quite different. In experiment 1, there are two observations for each person (caffeinated and decaf) and they are not statistically independent. In experiment 2, the data is independent.

If you have paired observations (experiment 1) and erroneously analyze them as unpaired observations (experiment 2), then you could get this result with a p -value of 0.9867:

```
> t.test(x,y)

Welch Two Sample t-test

data:  x and y
t = -0.0166, df = 198, p-value = 0.9867
alternative hypothesis: true difference in means is not equal to 0
95 percent confidence interval:
 -30.56737  30.05605
sample estimates:
mean of x mean of y
 501.2008  501.4565
```

The large p -value forces you to conclude that there is no difference between the groups. Contrast that result with the one from analyzing the same data, but correctly identifying it as paired:

```
> t.test(x, y, paired=TRUE)

Paired t-test

data:  x and y
t = -2.3636, df = 99, p-value = 0.02005
alternative hypothesis: true difference in means is not equal to 0
95 percent confidence interval:
 -0.4702824 -0.0410375
sample estimates:
mean of the differences
 -0.2556599
```

The p -value plummets to 0.02005, and we reach exactly the opposite conclusion.

See Also

If the populations are not normally distributed (bell-shaped) and either sample is small, the t test is inappropriate. In that case, consider using the Wilcoxon-Mann-Whitney test, which is implemented by the `wilcox.test` function.

1.15 Testing a Correlation for Significance

Problem

You calculated the correlation between two variables, but you don't know if the correlation is statistically significant.

Solution

The `cor.test` function can calculate both the p -value and the confidence interval of the correlation. If the variables came from normally distributed populations, use the default measure of correlation, known as the Pearson method:

```
> cor.test(x, y)
```

For nonnormal populations, use the Spearman method instead:

```
> cor.test(x, y, method="Spearman")
```

The function returns several values, including the p -value from the test of significance. Conventionally, $p < 0.05$ indicates that the correlation is likely significant, whereas $p > 0.05$ indicates it is not.

Discussion

In my experience, people often fail to check a correlation for significance. In fact, many people are unaware that a correlation can be insignificant. They jam their data into a computer, calculate the correlation, and blindly believe the result. However, they should ask themselves: Was there enough data? Is the magnitude of the correlation large enough? Fortunately, the `cor.test` function answers those questions.

Suppose we have two vectors, x and y , with values from normal populations. We might be very pleased that their correlation is greater than 0.83:

```
> cor(x, y)
[1] 0.8352458
```

But that is naïve. If we run `cor.test`, it reports a relatively large p -value of 0.1648:

```
> cor.test(x, y)

Pearson's product-moment correlation

data:  x and y
t = 2.1481, df = 2, p-value = 0.1648
alternative hypothesis: true correlation is not equal to 0
95 percent confidence interval:
 -0.6379590  0.9964437
sample estimates:
      cor
0.8352458
```

The p -value is above the conventional threshold of 0.05, so we conclude that the correlation is unlikely to be significant.

You can also check the correlation by using the confidence interval. In this example, the confidence interval is $(-0.638, 0.996)$. The interval contains zero, so it is possible that the correlation is zero, in which case there would be no correlation. Again, you could not be confident that the reported correlation is significant.

The `cor.test` output also includes the point estimate reported by `cor` (at the bottom, labeled “sample estimates”), saving you the additional step of running `cor`.

By default, `cor.test` calculates the Pearson correlation, which assumes that the underlying populations are normally distributed. The Spearman method makes no such assumption because it is nonparametric. Use `method="Spearman"` when working with nonnormal data.

See Also

See [Recipe 1.8](#) for computing correlations and other basic statistics.

1.16 Creating a Scatter Plot

Problem

You have paired observations: (x_1, y_1) , (x_2, y_2) , ..., (x_n, y_n) . You want to create a scatter plot of the pairs.

Solution

If your data is held in two parallel vectors, x and y , use them as arguments of `plot`:

```
> plot(x, y)
```

If your data is held in a (two-column) data frame, plot the data frame:

```
> plot(dfm)
```

Discussion

A scatter plot is usually my first attack on a new dataset. It's a quick way to see the relationship, if any, between x and y . Creating the scatter plot is easy:

```
> plot(x, y)
```

The `plot` function does not return anything. Rather, its purpose is to draw a plot of the (x, y) pairs in the graphics window.

Life is even easier if your data is captured in a two-column data frame. If you `plot` a two-column data frame, the function assumes you want a scatter plot created from those two columns. The scatter plot shown in [Figure 1-3](#) was created by one call to `plot`:

```
> plot(cars)
```

The `cars` dataset contains two columns: `speed` and `dist`. The first column is `speed`, so that becomes the x -axis, and `dist` becomes the y -axis.

If your data frame contains more than two columns, you will get multiple scatter plots, which might or might not be useful.

To get a scatter plot, your data must be numeric. Recall that `plot` is a polymorphic function, so, if the arguments are nonnumeric, it will create some other type of plot. See [Recipe 1.18](#), for example, which creates box plots from factors.

See Also

See the help page for `plot` to learn more about adding a title, subtitle, and labels.




Figure 1-3. Scatter plot

1.17 Creating a Bar Chart

Problem

You want to create a bar chart.

Solution

Use the `barplot` function. The first argument is a vector of bar heights:

```
> barplot(c(height1, height2, height3))
```

Discussion

The `barplot` function produces a simple bar chart. It assumes that the heights of your bars are conveniently stored in a vector. That is not always the case, however. Often, you have a vector of numeric data and a parallel factor that groups the data, and you want to produce a bar chart of the group means or the group totals. For example, the `airquality` dataset contains a numeric `Temp` column and a `Month` column. We can create

a bar chart of the mean temperature by month in two steps. First, we compute the means:

```
> heights <- tapply(airquality$Temp, airquality$Month, mean)
```

That gives the heights of the bars, from which we create the bar chart:

```
> barplot(heights)
```

The result is shown in the lefthand panel of [Figure 1-4](#). The result is pretty bland, as you can see, so it's common to add some simple adornments: a title, labels for the bars, and a label for the y-axis:

```
> barplot(heights,
+         main="Mean Temp. by Month",
+         names.arg=c("May", "Jun", "Jul", "Aug", "Sep"),
+         ylab="Temp (deg. F)")
```




Figure 1-4. Bar charts, plain (left) and adorned (right)

The righthand panel of [Figure 1-4](#) shows the improved bar chart.

See Also

The `barchart` function in the `lattice` package is another way to produce a bar chart.

1.18 Creating a Box Plot

Problem

You want to create a box plot of your data.

Solution

Use `boxplot(x)`, where `x` is a vector of numeric values.

Discussion

A box plot provides a quick and easy visual summary of a dataset. Figure 1-5 shows a typical box plot:

- The thick line in the middle is the median.
- The box surrounding the median identifies the first and third quartiles; the bottom of the box is Q1, and the top is Q3.
- The “whiskers” above and below the box show the range of the data, excluding outliers.
- The circles identify outliers. By default, an outlier is defined as any value that is farther than $1.5 \times \text{IQR}$ away from the box. (IQR is the *interquartile range*, or $Q3 - Q1$.) In this example, there are three outliers.




Figure 1-5. Your typical box plot

1.19 Creating a Histogram

Problem

You want to create a histogram of your data.

Solution

Use `hist(x)`, where `x` is a vector of numeric values.

Discussion

The lefthand panel of [Figure 1-6](#) shows a histogram of the `MPG.city` column taken from the `Cars93` dataset, created like this:

```
> data(Cars93, package="MASS")
> hist(Cars93$MPG.city)
```




Figure 1-6. Histograms

The `hist` function must decide how many cells (bins) to create for binning the data. In this example, the default algorithm chose seven bins. That creates too few bars for my taste, because the shape of the distribution remains hidden. So, I would include a second argument for `hist`—namely, the suggested number of bins:

```
> hist(Cars93$MPG.city, 20)
```

The number is only a suggestion, but `hist` will expand the number of bins as possible to accommodate that suggestion.

The righthand panel of [Figure 1-6](#) shows a histogram for the same data, but with more bins and with replacements for the default title and x-axis label. It was created like this:

```
> hist(Cars93$MPG.city, 20, main="City MPG (1993)", xlab="MPG")
```

See Also

The `histogram` function of the `lattice` package is an alternative to `hist`.

1.20 Performing Simple Linear Regression

Problem

You have two vectors, x and y , that hold paired observations: (x_1, y_1) , (x_2, y_2) , ..., (x_n, y_n) . You believe there is a linear relationship between x and y , and you want to create a regression model of the relationship.

Solution

The `lm` function performs a linear regression and reports the coefficients:

```
> lm(y ~ x)
```

```
Call:
```

```
lm(formula = y ~ x)
```

```
Coefficients:
```

```
(Intercept)          x
      17.72         3.25
```

Discussion

Simple linear regression involves two variables: a *predictor variable*, often called x ; and a *response variable*, often called y . The regression uses the ordinary least-squares (OLS) algorithm to fit the linear model:

$$y_i = \beta_0 + \beta_1 x_i + \varepsilon_i$$

where β_0 and β_1 are the regression coefficients and ε_i represents the error terms.

The `lm` function can perform linear regression. The main argument is a *model formula*, such as `y ~ x`. The formula has the response variable on the left of the tilde character (`~`) and the predictor variable on the right. The function estimates the regression coefficients, β_0 and β_1 , and reports them as the intercept and the coefficient of x , respectively:

```
Coefficients:
```

```
(Intercept)          x
      17.72         3.25
```

In this case, the regression equation is:

$$y_i = 17.72 + 3.25x_i + \varepsilon_i$$

It is quite common for data to be captured inside a data frame, in which case you want to perform a regression between two data frame columns. Here, *x* and *y* are columns of a data frame *dfrm*:

```
> dfrm
      x      y
1 0.04781401 5.406651
2 1.90857986 19.941568
3 2.79987246 23.922613
4 4.46755305 32.432904
5 3.76490363 44.259268
6 5.92364632 61.151480
7 8.04611587 26.305505
8 7.11097986 43.606087
9 9.73645966 58.262112
10 9.19324543 57.631029
.
. (etc.)
.
```

The `lm` function lets you specify a data frame by using the `data` parameter. If you do, the function will take the variables from the data frame and not from your workspace:

```
> lm(y ~ x, data=dfrm)      # Take x and y from dfrm
```

Call:

```
lm(formula = y ~ x, data = dfrm)
```

Coefficients:

```
(Intercept)      x
      17.72      3.25
```

1.21 Performing Multiple Linear Regression

Problem

You have several predictor variables (e.g., *u*, *v*, and *w*) and a response variable (*y*). You believe there is a linear relationship between the predictors and the response, and you want to perform a linear regression on the data.

Solution

Use the `lm` function. Specify the multiple predictors on the righthand side of the formula, separated by plus signs (+):

```
> lm(y ~ u + v + w)
```

Discussion

Multiple linear regression is the obvious generalization of simple linear regression. It allows multiple predictor variables instead of one predictor variable and still uses OLS to compute the coefficients of a linear equation. The three-variable regression just shown corresponds to this linear model:

$$y_i = \beta_0 + \beta_1 u_i + \beta_2 v_i + \beta_3 w_i + \varepsilon_i$$

R uses the `lm` function for both simple and multiple linear regression. You simply add more variables to the righthand side of the model formula. The output then shows the coefficients of the fitted model:

```
> lm(y ~ u + v + w)

Call:
lm(formula = y ~ u + v + w)

Coefficients:
(Intercept)          u          v          w
    1.4222      1.0359      0.9217      0.7261
```

The `data` parameter of `lm` is especially valuable when the number of variables increases, since it's much easier to keep your data in one data frame than in separate variables. Suppose your data is captured in a data frame, such as the `dfrm` variable shown here:

```
> dfrm
   y      u      v      w
1  6.584519 0.79939065 2.7971413 4.366557
2  6.425215 -2.31338537 2.7836201 4.515084
3  7.830578 1.71736899 2.7570401 3.865557
4  2.757777 1.27652888 0.4191765 2.547935
5  5.794566 0.39643488 2.3785468 3.265971
6  7.314611 1.82247760 1.8291302 4.518522
7  2.533638 -1.34186107 2.3472593 2.570884
8  8.696910 0.75946803 3.4028180 4.442560
9  6.304464 0.92000133 2.0654513 2.835248
10 8.095094 1.02341093 2.6729252 3.868573
.
. (etc.)
.
```

When we supply `dfrm` to the `data` parameter of `lm`, R looks for the regression variables in the columns of the data frame:

```
> lm(y ~ u + v + w, data=dfrm)

Call:
lm(formula = y ~ u + v + w, data = dfrm)

Coefficients:
(Intercept)          u          v          w
    1.4222      1.0359      0.9217      0.7261
```


See Also

See [Recipe 1.20](#) for simple linear regression.

1.22 Getting Regression Statistics

Problem

You want the critical statistics and information regarding your regression, such as R^2 , the F statistic, confidence intervals for the coefficients, residuals, the ANOVA table, and so forth.

Solution

Save the regression model in a variable, say `m`:

```
> m <- lm(y ~ u + v + w)
```

Then, use the following functions to extract regression statistics and information from the model:

`anova(m)`

ANOVA table

`coefficients(m)`

Model coefficients

`coef(m)`

Same as `coefficients(m)`

`confint(m)`

Confidence intervals for the regression coefficients

`deviance(m)`

Residual sum of squares

`effects(m)`

Vector of orthogonal effects

`fitted(m)`

Vector of fitted y values

`residuals(m)`

Model residuals

`resid(m)`

Same as `residuals(m)`

`summary(m)`

Key statistics, such as R^2 , the F statistic, and the residual standard error (σ)

`vcov(m)`

Variance-covariance matrix of the main parameters

Discussion

When I started using R, the documentation said to use the `lm` function to perform linear regression. So I did something like this, getting the output shown in [Recipe 1.21](#):

```
> lm(y ~ u + v + w)

Call:
lm(formula = y ~ u + v + w)

Coefficients:
(Intercept)          u          v          w
    1.4222     1.0359     0.9217     0.7261
```

I was so disappointed! The output was nothing compared to other statistics packages such as SAS. Where is R^2 ? Where are the confidence intervals for the coefficients? Where is the F statistic, its p -value, and the ANOVA table?

Of course, all that information is available—you just have to ask for it. Other statistics systems dump everything and let you wade through it. R is more minimalist. It prints a bare-bones output and lets you request what more you want.

The `lm` function returns a *model object*. You can save the object in a variable by using the assignment operator (`<-`). This example assigns the object to the variable `m`:

```
> m <- lm(y ~ u + v + w)
```

From the model object, you can extract important information using specialized functions. The most important function is `summary`:

```
> summary(m)

Call:
lm(formula = y ~ u + v + w)

Residuals:
    Min       1Q   Median       3Q      Max
-3.3965 -0.9472 -0.4708  1.3730  3.1283

Coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept)   1.4222     1.4036   1.013  0.32029
u             1.0359     0.2811   3.685  0.00106 **
v             0.9217     0.3787   2.434  0.02211 *
w             0.7261     0.3652   1.988  0.05744 .
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.625 on 26 degrees of freedom
Multiple R-squared:  0.4981,    Adjusted R-squared:  0.4402
F-statistic: 8.603 on 3 and 26 DF,  p-value: 0.0003915
```

The summary shows the estimated coefficients. It shows the critical statistics, such as R^2 , and the F statistic. It also shows an estimate of σ , the standard error of the residuals.

There are also specialized extractor functions for other important information:

Model coefficients (point estimates)

```
> coef(m)
(Intercept)          u          v          w
  1.4222050   1.0358725   0.9217432   0.7260653
```

Confidence intervals for model coefficients

```
> confint(m)
                2.5 %   97.5 %
(Intercept) -1.46302727  4.307437
u             0.45805053  1.613694
v             0.14332834  1.700158
w            -0.02466125  1.476792
```

Model residuals

```
> resid(m)
 1          2          3          4          5          6
-1.41440465  1.55535335 -0.71853222 -2.22308948 -0.60201283 -0.96217874
 7          8          9         10         11         12
-1.52877080  0.12587924 -0.03313637  0.34017869  1.28200521 -0.90242817
13          14         15         16         17         18
 2.04481731  1.13630451 -1.19766679 -0.60210494  1.79964497  1.25941264
19          20         21         22         23         24
-2.03323530  1.40337142 -1.25605632 -0.84860707 -0.47307439 -0.76335244
25          26         27         28         29         30
 2.16275214  1.53483492  1.65085364 -3.39647629 -0.46853750  3.12825629
```

Residual sum of squares

```
> deviance(m)
[1] 68.69616
```

ANOVA table

```
> anova(m)
Analysis of Variance Table

Response: y
  Df Sum Sq Mean Sq F value    Pr(>F)
u   1  27.916  27.9165  10.5658 0.003178 **
v   1  29.830  29.8299  11.2900 0.002416 **
w   1  10.442  10.4423   3.9522 0.057436 .
Residuals 26  68.696   2.6422
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

If you find it annoying to save the model in a variable, you are welcome to use one-liners such as this:

```
> summary(lm(y ~ u + v + w))
```

1.23 Diagnosing a Linear Regression

Problem

You have performed a linear regression. Now you want to verify the model's quality by running diagnostic checks.

Solution

Start by plotting the model object, which will produce several diagnostic plots:

```
> m <- lm(y ~ x)
> plot(m)
```

Next, identify possible outliers either by looking at the diagnostic plot of the residuals or by using the `outlier.test` function of the `car` package:

```
> library(car)
> outlier.test(m)
```

Finally, identify any overly influential observations (by using the `influence.measures` function, for example).

Discussion

R fosters the impression that linear regression is easy: just use the `lm` function. Yet fitting the data is only the beginning. It's your job to decide whether the fitted model actually works and works well.

Before anything else, you must have a statistically significant model. Check the F statistic from the model summary ([Recipe 1.22](#)) and be sure that the p -value is small enough for your purposes. Conventionally, it should be less than 0.05, or else your model is likely meaningless.

Simply plotting the model object produces several useful diagnostic plots:

```
> m <- lm(y ~ x)
> plot(m)
```

[Figure 1-7](#) shows diagnostic plots for a pretty good regression:

- The points in the Residuals vs Fitted plot are randomly scattered with no particular pattern.
- The points in the Normal Q–Q plot are more-or-less on the line, indicating that the residuals follow a normal distribution.
- In both the Scale–Location plot and the Residuals vs Leverage plots, the points are in a group with none too far from the center.




Figure 1-7. Diagnostic plots: pretty good fit

In contrast, [Figure 1-8](#) shows the diagnostics for a not-so-good regression. Observe that the Residuals vs Fitted plot has a definite parabolic shape. This tells us that the model is incomplete: a quadratic factor is missing that could explain more variation in y . Other patterns in residuals are suggestive of additional problems; a cone shape, for example, may indicate nonconstant variance in y . Interpreting those patterns is a bit of an art, so I suggest reviewing a good book on linear regression while evaluating the plot of residuals.

There are other problems with the not-so-good diagnostics. The Normal Q-Q plot has more points off the line than it does for the good regression. Both the Scale-Location and Residuals vs Leverage plots show points scattered away from the center, which suggests that some points have excessive leverage.




Figure 1-8. Diagnostic plots: not-so-good fit

Another pattern is that point number 28 sticks out in every plot. This warns us that something is odd with that observation. The point could be an outlier, for example. We can check that hunch with the `outlier.test` function of the `car` package:

```
> outlier.test(m)

max|tstudent| = 3.183304, degrees of freedom = 27,
unadjusted p = 0.003648903, Bonferroni p = 0.1094671
```

Observation: 28

`outlier.test` identifies the model's most outlying observation. In this case, it identified observation number 28 and so confirmed that it could be an outlier.

See Also

The `car` package is not part of the standard distribution of R; download and install it using the `install.packages` function.

1.24 Predicting New Values

Problem

You want to predict new values from your regression model.

Solution

Save the predictor data in a data frame. Use the `predict` function, setting the `newdata` parameter to the data frame:

```
> m <- lm(y ~ u + v + w)
> preds <- data.frame(u=3.1, v=4.0, w=5.5)
> predict(m, newdata=preds)
```

Discussion

Once you have a linear model, making predictions is quite easy because the `predict` function does all the heavy lifting. The only annoyance is arranging for a data frame to contain your data.

The `predict` function returns a vector of predicted values with one prediction for every row in the data. The example in the Solution contains one row, so `predict` returns one value:

```
> preds <- data.frame(u=3.1, v=4.0, w=5.5)
> predict(m, newdata=preds)
      1
12.31374
```

If your predictor data contains several rows, you get one prediction per row:

```
> preds <- data.frame(
+   u=c(3.0, 3.1, 3.2, 3.3),
+   v=c(3.9, 4.0, 4.1, 4.2),
+   w=c(5.3, 5.5, 5.7, 5.9) )
> predict(m, newdata=preds)
      1      2      3      4
11.97277 12.31374 12.65472 12.99569
```

In case it's not obvious, the new data needn't contain values for response variables, only predictor variables. After all, you are trying to *calculate* the response, so it would be unreasonable of R to expect you to supply it.

See Also

These are just the point estimates of the predictions. Use the `interval="prediction"` argument of `predict` to obtain the confidence intervals.

1.25 Accessing the Functions in a Package

Problem

A package installed on your computer is either a standard package or a package downloaded by you. When you try using functions in the package, however, R cannot find them.

Solution

Use either the `library` function or the `require` function to load the package into R:

```
> library(packagename)
```

Discussion

R comes with several standard packages, but not all of them are automatically loaded when you start R. Likewise, you can download and install many useful packages from CRAN, but they are not automatically loaded when you run R. The `MASS` package comes standard with R, for example, but you could get this message when using the `lda` function in that package:

```
> lda(x)
Error: could not find function "lda"
```

R is complaining that it cannot find the `lda` function among the packages currently loaded into memory.

When you use the `library` function or the `require` function, R loads the package into memory and its contents immediately become available to you:

```
> lda(f ~ x + y)
Error: could not find function "lda"
> library(MASS)                # Load the MASS library into memory
> lda(f ~ x + y)                # Now R can find the function
Call:
lda(f ~ x + y)

Prior probabilities of groups:
.
. (etc.)
.
```

Before calling `library`, R does not recognize the function name. Afterward, the package contents are available and calling the `lda` function works.

Notice that you needn't enclose the package name in quotes.

The `require` function is nearly identical to `library`, but it has two features that are useful for writing scripts. It returns `TRUE` if the package was successfully loaded and `FALSE` otherwise. It also generates a mere warning if the load fails—unlike `library`, which generates an error.

Both functions have a key feature: they do not reload packages that are already loaded, so calling twice for the same package is harmless. This is especially nice for writing scripts. You can write a script to load needed packages while knowing that loaded packages will not be reloaded:

The `detach` function will unload a package that is currently loaded.

```
> detach(package:MASS)
```

Observe that the package name must be qualified, as in `package:MASS`.

One reason to unload a package is if it contains a function whose name conflicts with a same-named function lower on the search list. When such a conflict occurs, we say the higher function *masks* the lower function. You no longer “see” the lower function because R stops searching when it finds the higher function. Hence unloading the higher package unmasks the lower name.

See Also

See the `search` function for more about the search path.